Source-Level Reliability Estimation under Publisher-Disjoint Evaluation

Anonymous Author(s)

Affiliation withheld for double-blind review

Abstract. Whether a news outlet can be trusted is conventionally judged by expert organisations, which weigh editorial practice, factual accuracy and transparency at the level of the outlet itself. Those judgements are careful but manual, and the outlets they never reach stay unrated. Computational formulations of the same task usually operate one article at a time: they fit a classifier to pre-labelled items and score it on held-out ones. We instead cast the task as prediction about the source. Our framework has two stages: transformer-based language models rate individual articles, and we then pool those ratings to characterise outlets the models never saw during training. To keep the setting close to deployment, every split is publisher-disjoint, so no outlet contributes to both the training and the test side. We work with 19,476 English-language political news items from 439 outlets carrying NewsGuard ratings. Exact accuracy rises from roughly 0.60 on single articles to 0.69 on outlets, and the vote itself accounts for one to four of those points. Counting a prediction one level away from the true rating as correct, accuracy rises from 0.84 to 0.88. The labels invite that tolerance, because the five levels bin a continuous 0–100 score. We then ask how the errors distribute across the political spectrum, and find systematic associations between the leaning of an outlet and the mistakes made on its articles. We report that disparity as a bound on responsible use. Taken together, the results indicate that text alone carries enough signal to place an unseen outlet on a reliability scale, and we set out the conditions under which such estimates may direct expert attention rather than replace it.

Keywords: News source reliability · Publisher reliability inference · Transformer-based text classification · Cross-publisher generalization · Reliability aggregation · Political bias analysis.

Abstract
(1) R1.2: il salto 0.60 -> 0.69 confondeva l'aggregazione con il cambio di unita'. Misurato sugli stessi 439 publisher, il voto vale da +1.6 a +4.3 punti, il resto e' il passaggio da articoli a testate. Riformulato il claim. (2) L'abstract e' parafrasato rispetto alla versione su arXiv: serviva a non far risalire a quel PDF da un motore di ricerca, cosa che avrebbe compromesso il doppio cieco. (3) Aggiunti anche i risultati tolerant, solo accennati.

1 Introduction

Assessing the reliability of news publishers has become increasingly important for limiting the spread of misleading or low-quality information. Professional organizations such as NewsGuard, the Global Disinformation Index, and Media Bias Fact Check evaluate news outlets according to journalistic standards including transparency, editorial accountability, factual accuracy, and correction

practices, and despite methodological differences they exhibit substantial convergence [11,19]. Importantly, these assessments are performed at the level of the publisher rather than of individual articles. Trained analysts produce them by hand, so the outlets they never reach stay unrated.

A large body of computational work instead formulates the problem as article-level classification [1,14,21,24]. This only partially reflects the structure of professional assessment, where reliability is a persistent property of the source emerging across multiple publications: highly reliable outlets occasionally publish lower-quality content, and unreliable publishers produce articles that appear plausible in isolation.

A second limitation concerns evaluation. Existing studies derive article labels from publisher-level annotations without enforcing publisher-disjoint splits [1,6, 23], so models can partially memorize publisher-specific patterns rather than learning transferable indicators of reliability [3,26].

We address these limitations by reformulating the task as publisher-level inference. We propose a two-stage framework in which transformer-based language models first estimate the reliability of individual articles and then aggregate those predictions to infer the reliability of previously unseen publishers, under a strict publisher-disjoint evaluation protocol. Our contributions are as follows: we reformulate news reliability assessment as a source-level inference problem; we introduce a publisher-disjoint protocol that prevents source leakage and measures generalization to unseen publishers; we construct a dataset of 19,476 political news articles from 439 English-language publishers labeled with NewsGuard scores; and we show that aggregating article-level predictions improves reliability estimation, reaching 0.69 publisher-level accuracy on unseen sources. An estimate of this kind carries consequences for the outlets it ranks, so we also audit how the errors distribute across the political spectrum. Section 6 states the societal need we address, the preconditions any benefit depends on, and the harms the approach may introduce.

2 Related Work

Early work on online news quality focused on article-level classification, detecting misleading content through linguistic features or neural architectures [14], often without considering the publisher.

Publisher reliability and the signals used to estimate it A parallel line shifts the focus to the source, variously termed *reliability*, *trustworthiness* or *credibility* [8, 15,27], whose reputation is known to influence perceived information quality [9]. Computational approaches rely on textual and stylistic features [21] or on external signals such as user interactions, social networks and metadata [4,13,20,28]; journalistic rating systems are commonly used as ground truth [11,19], and recent work explores LLM-based estimation [10,18,24,29].

Datasets and cross-domain generalization Datasets with explicit publisher reliability labels remain limited [1,5,23], and recent work shows that performance degrades under distribution shifts across domains and publishers [3,6,26].

Our contribution In contrast, we focus exclusively on textual content, enforce a strict publisher-disjoint evaluation protocol and report the error disparities the resulting estimates carry.

3 Problem Definition and Methodology

3.1 Problem definition

We formulate the problem as 'News Source Reliability Inference': estimating the reliability level of a publisher from its published content. Because publishers are observed through collections of articles, we decompose it into (i) article-level reliability estimation and (ii) publisher-level inference through aggregation.

Let A be a corpus of news articles, each with a text in $\mathcal{T}$ and a publisher $p(a) \in P$, and let $A_p = \{a \in A \mid p(a) = p\}$ be the articles of p. Each publisher has a fixed reliability level $r_P(p)$ over five ordered categories $L = \{l_1, \ldots, l_5\}$ obtained by binning continuous scores (Section 4); since reliability is a publisher's property, every article inherits it, $r_A(a) = r_P(p(a))$. A political leaning $o_P(p) \in O$, with O ranging over Far Left, Slightly Left, Non labelled, Slightly Right and Far Right, is inherited in the same way. Both inheritance functions are needed because article-level annotation is unavailable; political leaning is never a model input or target, and is retained only for the post-hoc analysis of Section 5.4.

In **Stage 1** we train a classifier $C_{\text{reliability}} : \mathcal{T} \to L$ to predict the reliability of a single article from its text alone. In **Stage 2** we aggregate those predictions through $Agg : L^{|A_p|} \to L$, so that $\hat{r}_P(p) = Agg(\{C_{\text{reliability}}(a) \mid a \in A_p\})$; Section 5.5 instantiates Agg as the mode, i.e. majority voting.

Because each article inherits its publisher's label, articles from the same source are not independently labeled, and a random split at the article level would leak test labels into training. With P_{train} and P_{test} the publishers of the training and test articles, we therefore enforce

$$P_{\text{train}} \cap P_{\text{test}} = \emptyset. \tag{1}$$

All test articles thus originate from previously unseen publishers, which aligns the evaluation with the source-level nature of the labels and prevents publisher-specific memorization. For the same reason we mask no source identifier. A publisher-specific cue cannot help at test time, because the model never sees a test publisher during training.

R1.5 Chiedeva masking degli identificatori di publisher e ablation su testo sanificato. Risposto nel testo: con lo split publisher-disjoint un indizio specifico del publisher e' inutilizzabile in test. Aggiunta una frase, nessun esperimento nuovo.

3.2 Evaluation Metrics and Baselines

We use the same metrics for both article-level and source-level evaluations.

Standard and Macro Metrics We report Accuracy as the ratio of correct predictions. Because class frequencies differ moderately (l_4 is about 2.6 times larger than l_5), we prefer Macro-averaged Precision, Recall and F1, which treat all five classes equally.

Tolerant Metrics Given the ordinal nature of our labels, standard accuracy may be overly punitive for minor deviations. We therefore adopt *Tolerant Accuracy* [17], which assigns partial credit to predictions falling into adjacent categories. Let N be the total number of samples and c_{ij} the count of instances with true label i predicted as j:

$$\text{Accuracy}_{\text{tolerant}} = \frac{1}{N} \sum_{i \in L} \sum_{j \in L} w_{ij} \cdot c_{ij}, \qquad w_{ij} = \begin{cases} 1 & \text{if } |i-j| = 0 \\ \alpha & \text{if } |i-j| = 1 \\ 0 & \text{if } |i-j| > 1 \end{cases} \tag{2}$$

The parameter $\alpha \in [0, 1]$ controls the tolerance level. We report results for $\alpha = 0.5$ (half credit for adjacent errors) and $\alpha = 1.0$ (full credit, effectively treating adjacent classes as correct).

Baselines We pair each baseline with the metrics it naturally challenges. The **Majority Baseline** assigns every instance to the most frequent class observed in the training set, and achieves the maximum accuracy obtainable without processing the input text. The **Uniform Random Baseline** assigns labels uniformly at random, giving each of the five reliability levels an equal probability of 20%. We use it for Macro metrics, for which a Majority classifier would be an unsuitable benchmark, as it would score zero Recall and F1 on all other categories, and it establishes a floor of 0.20.

3.3 Bias Assessment and Statistical Analysis

Beyond aggregate performance, we investigate *how* and *why* models fail.

Error Analysis We compute the Error Rate per level, $(1 - \text{Recall}_l)$. To identify systematic over- or under-estimation we analyze the Mean Error Direction (MED) over misclassified samples:

$$\text{MED} = \frac{1}{|E|} \sum_{i \in E} (\hat{y}_i - y_i) \tag{3}$$

where $\hat{y}_i$ and y_i represent the predicted and true reliability levels (mapped to their ordinal values $1, \ldots, 5$), respectively. Positive values indicate a tendency to overestimate reliability, while negative values indicate a tendency to underestimate reliability.

Statistical Significance of Political Bias To determine whether classification errors depend on the publisher's political leaning, we conduct a Pearson's χ^2 test of independence. We test the null hypothesis (H_0) that error frequency is independent of political orientation.

To quantify the magnitude of this effect independent of sample size, we compute Cramér's V:

$$V = \sqrt{\frac{\chi^2}{N \cdot \min(r-1, c-1)}} \tag{4}$$

where N is the total number of samples, while r and c represent the number of rows (political orientations) and columns (classification outcomes) in the contingency table.

4 Dataset and Labeling

Ground Truth: NewsGuard The publisher-level reliability function $r_P : P \to L$ is derived from NewsGuard,[1] a professional media rating system in which trained journalists evaluate online news outlets. Each publisher $p \in P$ is assessed against nine widely adopted journalistic criteria, including deceptive headlines, transparency of ownership and financing, and correction practices.

The evaluation produces a continuous reliability score in $[0, 100]$, measuring a publishers adherence to professional journalistic standards.[2] The proprietary ratings were licensed for use to the authors of this publication.

Label Construction and Publishers Selection We discretize the continuous score into five ordinal categories $L = \{l_1, \ldots, l_5\}$ following NewsGuard's own thresholds. Table 1 reports the mapping, with the number of publishers per label after filtering.

Label	Score Range	Initial Dom.	Scraped Dom.	Final Art.
l_1	0–39 (Max Caution)	~~1,211~~ → 175	152	4,871
l_2	40–59 (Caution)	127	69	3,016
l_3	60–74 (Credible w/ Exc.)	124	95	4,620
l_4	75–99 (Generally Credible)	175	88	5,041
l_5	100 (High Credibility)	68	35	1,928

Table 1. Reliability labels and NewsGuard score ranges, with the number of domains after filtering, successfully scraped domains, and collected articles.

We restrict P to outlets tagged as *Political News and Commentary*, publishing mainly in English and without a paywall. To prevent the corpus from being

[1] https://www.newsguardtech.com/

[2] https://www.newsguardtech.com/ratings/rating-process-criteria/

dominated by highly unreliable sources, we cap l_1 at 175 publishers, matching l_4; the final set contains $|P| = 669$ outlets.

Article Collection Let A denote the set of collected news articles. For each publisher $p \in P$, we scrape articles from its homepage using the `newspaper3k` library[3], collecting at most 200 articles per outlet to ensure a balanced representation across publishers. The scraper is configured to retrieve English-language content only and to prevent duplicates. Scraping did not succeed for all outlets, owing to anti-scraping measures, timeouts or structural incompatibilities, so the final publisher count is lower than the original selection. We retain only articles of at least 100 words, excluding short announcements and advertisements. This procedure yields an initial corpus of 25,822 articles.

Data Curation and Labeling To preserve the original linguistic characteristics of the articles, we apply no text normalization or cleaning. Some domains nonetheless returned localized content, which we filter out with the `langdetect` library[4]. This process removes 6,346 articles and leaves 19,476 English-language articles in the dataset.

The final composition is given in the last two columns of Table 1; l_5 remains the least represented class, reflecting its limited presence in the original NewsGuard data.

Political Leaning Annotation NewsGuard also provides a political leaning annotation for a subset of publishers, inherited at the article level as $o_A(a) = o_P(p(a))$. Reliability levels are unevenly represented across political categories, a structural property of the media ecosystem that must be considered when interpreting model performance and error patterns in the analyses that follow.

5 Experiments

This section presents the setup and the results for the tasks defined in Section 3.1: (i) *Article-level reliability estimation* (Stage 1) under a publisher-disjoint protocol, and (ii) *Publisher reliability inference* (Stage 2) obtained by aggregating article-level predictions.

5.1 Experimental Setup

Following the validation protocol described in Section 3 (Equation 1), we use a 5-fold cross-validation scheme with strict separation of publishers between training and test sets. This assesses generalization to articles from previously unseen publishers and approximates a realistic *in-the-wild* deployment. Folds are stratified on publisher reliability, so that each preserves the same reliability distribution.

[3] https://pypi.org/project/newspaper3k/
[4] https://pypi.org/project/langdetect/

We select four transformer-based models spanning encoder baselines and decoder-based large language models. **BERT**-base-uncased [7] is our primary baseline, establishing a lower bound comparable with a vast body of prior work [22]. **RoBERTa**-large [12] tests whether optimized training strategies alone improve reliability detection with the architecture unchanged. **Llama 3** 8B [25] represents the shift toward decoder-only generative models, and **Mistral** 24B [16], roughly three times larger, tests whether additional capacity improves detection or whether performance saturates.

Training Strategies and Hyperparameters Encoders (BERT, RoBERTa) were *fully fine-tuned* with AdamW, a learning rate of $2e^{-5}$ and weight decay [7,12]. For the decoders (Llama, Mistral) full fine-tuning was computationally infeasible, so we adopted Low-Rank Adaptation (LoRA) with a learning rate of $2e^{-4}$ in mixed precision.

Each model was trained for 5 epochs per fold, except Mistral, which was trained for 3. This follows a preliminary convergence study (see Supplementary Material [2], which also reports the remaining optimisation and batching details) showing that validation loss minimizes at Epoch 4 for Llama, while Mistral converges within three epochs. We used early stopping on validation loss and a fixed global random seed (42). The maximum sequence length for input processing was set to 256 tokens, except for Mistral, for which it was increased to 1,024 tokens. This tests whether a longer context improves classification.

All fine-tuned models and code will be released to foster further research; they are withheld during review to preserve anonymity and will be made available upon acceptance.

Spostato
Nei supplementary: gli iperparametri completi (ottimizzatore, batching) e lo studio di convergenza con le curve di loss di Llama e Mistral.

Modelli
Non rilasciati durante la review: OSF non consente piu' di scaricare file di grandi dimensioni in modalita' anonima, quindi i pesi restano fuori e saranno pubblicati dopo l'accettazione, quando cade l'anonimato.

5.2 Stage 1: Article-level reliability estimation

We first evaluate the article-level estimators of Stage 1.

Figure 1 illustrates the results against the Majority Class Baseline (MCB), used for Accuracy and Tolerant Accuracy, and the Random Baseline (RB) at $1/N$, used for the macro metrics. All models reach a standard accuracy around 0.60, more than double the MCB, with Mistral highest at 0.61. On macro-averaged performance RoBERTa leads with $F1_{macro}$ of 0.55, followed by Mistral and Llama at 0.54 and BERT at 0.52, all well above the RB of 0.20; the variance across folds is minimal (standard deviation about 0.02). Under the stricter tolerance ($\alpha = 0.5$), Mistral, RoBERTa and Llama reach approximately 0.72 against an adjusted MCB of 0.43. With $\alpha = 1.0$ RoBERTa reaches 0.84 and the others 0.82–0.83; the baseline is naturally higher here (0.60), so the gain is +0.24. Misclassifications are therefore rarely severe, typically involving adjacent labels.

5.3 Biases in Article-level reliability estimation

Here, we study whether the classifier systematically performs better on some reliability classes than on others.

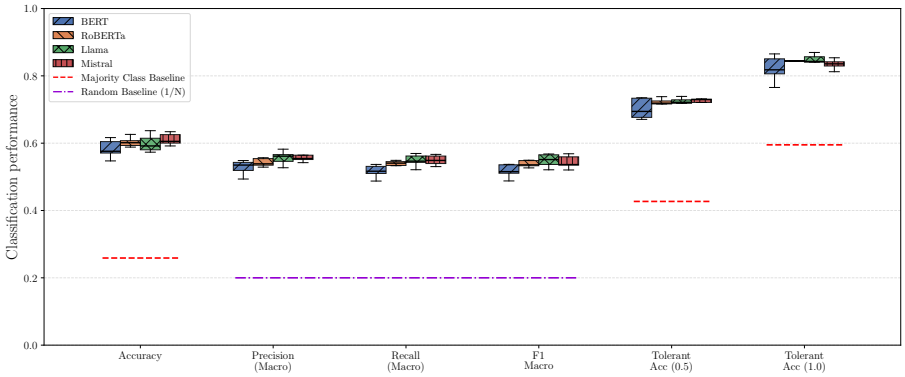


Fig. 1. Article-level metrics across the five folds. Horizontal lines are the baselines: dashed red for Majority Class, dash-dotted violet for Random.

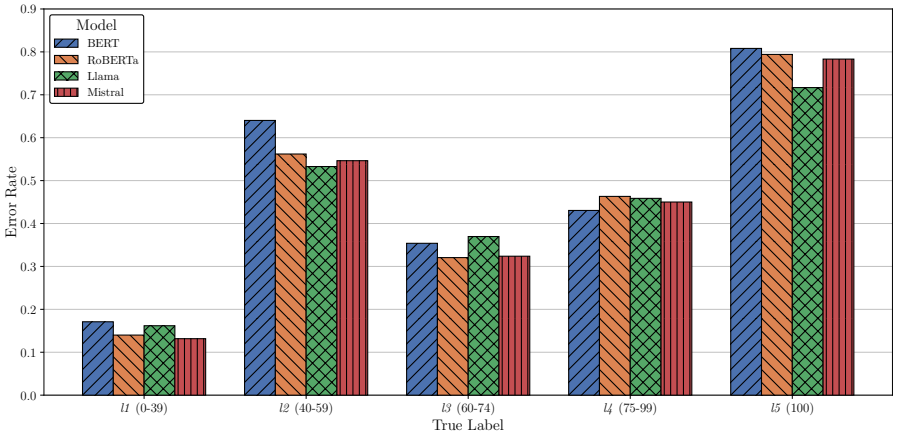


Fig. 2. Per-class error rate, $(1 - \text{Recall}_l)$.

Figure 2 highlights an asymmetry in prediction correctness. The lowest error rates are on l_1 (< 0.20), suggesting that highly unreliable content has recognizable structural features. The highest-reliability articles (l_5) are the most challenging to recognize for the model, with error rates exceeding 70%[5]. Class l_2 also proves difficult, with an error rate of over 50%, while l_3 sits in the middle at approximately 35%.

These patterns remain consistent across all four models, suggesting that the biases are structural properties of the task rather than of a specific architecture;

[5] Class l_5 only includes articles with a reliability score of 100, and is easily confused with l_4, which contains similarly high scores.

Model	Accuracy	χ^2	p-value	Cramérs V
BERT	58.27%	$1684.13 < 10^{-300}$		0.294
RoBERTa	60.34%	$1418.31 < 10^{-300}$		0.270
Llama	59.97%	$1207.83 < 10^{-250}$		0.249
Mistral	61.16%	$1343.98 < 10^{-300}$		0.263

Table 2. Political leaning vs. classification correctness: χ^2 test of independence with p-values, and Cramér's V as effect size.

the Supplementary Material [2] reports the error-direction analysis and the full confusion matrices.

Spostato
Nei supplementary: la figura del MED per livello di affidabilita' e le matrici di confusione a livello di articolo, per fold e in media.

5.4 Influence of Political Leaning on Classification Errors

We investigate whether the models treat all sides of the political spectrum equally. Political orientation is a property of the publisher, inherited by each article it produces (see Section 3.1), which enables a post-hoc analysis of how source-level alignment relates to article-level errors.

Error rates are computed separately for each leaning, which makes them independent of class frequency.

All models exhibit systematic variation across political leanings: there are substantially higher error rates for far left and far right outlets (5972%) compared to lower rates for slightly right and non-labelled outlets (3034%), and intermediate error rates for slightly left articles (4758%).

A chi-square test of independence rejects, for every model, the hypothesis that correctness is independent of political leaning (Table 2), with Cramér's V indicating a consistent small-to-moderate effect size. Leaning is a publisher-level attribute, and every article of an outlet carries the same value. We therefore repeat the test with one observation per publisher, and the association holds: χ^2 ranges from 18.8 to 44.8, $p < 10^{-3}$, and V from 0.21 to 0.32. Nor does reliability explain it away. We fit a logistic regression of error on leaning, control for the true reliability level and cluster the standard errors on the publisher. Far-left and far-right outlets still carry 4.0 to 6.5 times the odds of error of non-labelled ones, with $p < 0.05$ for every model (Supplementary Material [2]).

We then focus on misclassified instances to assess whether models systematically overestimate or underestimate reliability as a function of political leaning. Figure 3 reports the Mean Error Direction (MED) for each orientation (see Section 3.3); positive values indicate overestimation of reliability, negative values underestimation.

Clear directional patterns emerge across all four models. Far Left and Far Right articles exhibit consistent overestimation, with mean signed errors ranging from +0.33 to +1.00, approaching a full reliability level. Slightly Left articles show systematic underestimation, with average values between -0.54 and 0.82. Slightly Right and non-labelled outlets display smaller, less consistent deviations; the Supplementary Material [2] reports the corpus composition by leaning, the

R1.4 + R1.3
R1.4: il chi-quadro trattava 77.904 articoli come indipendenti. Rifatto a livello publisher (439 osservazioni): l'effetto resta, V sale a 0.21–0.32. R1.3: possibile confounding con il livello di affidabilita'. Logit con controllo per reliability e SE clusterizzati: OR 4.0–6.5 agli estremi. Aggiunte due frasi qui, tabelle complete nei supplementary.

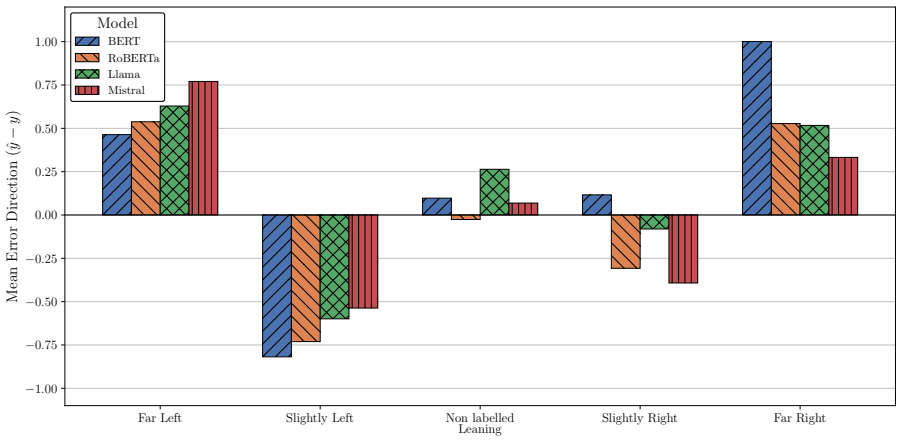


Fig. 3. Mean value and direction of misclassifications, stratified by political leaning. Positive values indicate overestimation (prediction > true label), while negative values indicate underestimation.

per-leaning error rates and the full confusion matrices stratified by political leaning.

Spostato
Nei supplementary: le figure della composizione del corpus per orientamento e dei tassi di errore per orientamento, piu' le matrici stratificate per leaning. Tolto anche il dettaglio numerico su Slightly Right e non etichettati.

5.5 Stage 2: Publisher reliability inference

Publisher-level inference is the primary objective of this work; the article-level estimators are intermediate components. We instantiate *Agg* as the mode of a publisher's article-level predictions, i.e. majority voting, assigning the most frequent class as $\hat{r}_P(p)$, on the hypothesis that combining predictions offsets the noise of individual classifications. Predictions are pooled across the five test folds, so each publisher's reliability is inferred only from articles the model never saw in training.

Figure 4 presents the results against the same two baselines used at the article level; the majority-class reference differs across the three accuracy variants, since a constant predictor also collects partial and then full credit under tolerance.

With an Accuracy of 0.69 RoBERTa leads, followed closely by Mistral at 0.68, more than triple the Random Baseline. A publisher-level cluster bootstrap nonetheless puts the four models within overlapping 95% intervals: RoBERTa 0.690 [0.647, 0.731], Mistral 0.683 [0.642, 0.727], BERT and Llama 0.665 [0.62, 0.71]. The ordering among them is therefore not itself a result. Considering $\text{Accuracy}_{\text{tolerant}}$ (1.0), Llama achieves the top score of 0.88, followed by RoBERTa (0.87) and Mistral (0.86), substantially outperforming the majority class baseline.

R1.8
Chiedeva intervalli di confidenza o test appaiati sulle differenze 0.69 vs 0.68. Aggiunto bootstrap cluster sui publisher: gli IC si sovrappongono, quindi il testo ora dice esplicitamente che la classifica fra modelli non e' un risultato. Tabella completa nei supplementary.

Publisher-level errors are strictly local, concentrating between l_4 (Generally Credible) and l_5 (High Credibility); the Supplementary Material [2] reports full publisher-level confusion matrices and a detailed error analysis.

Supplementary
Ospitati su Anonymous GitHub (anonymous.4open.science), che e' il link citato qui. Li carico anche nel progetto Overleaf, ma la copia che conta e' quella del repo: ogni modifica va fatta via GitHub, da dove si propaga al link anonimo.

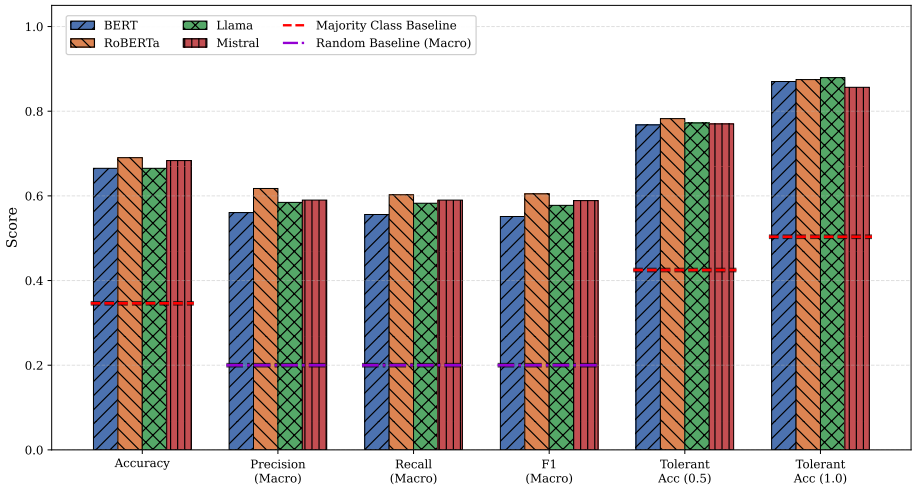


Fig. 4. Publisher-level metrics. Dashed red: Majority Class Baseline; dash-dotted violet: Random Baseline at 0.20.

6 Theory of Change

Work on misinformation is not socially beneficial by default. We therefore state what we expect this work to change, what must hold for it to do so, and how it may cause harm.

Societal Need and Intended Contribution The binding constraint on expert rating is the cost of *triage* rather than the quality of the judgement. We show that article text alone places an unseen outlet on a five-level scale with 0.69 accuracy (Section 5.2). Errors are predominantly local: a mistake usually lands on an adjacent level, not on a confusion between a disinformation site and an established newspaper. We claim no improvement in information quality itself, and we aim only to allocate expert effort.

Preconditions and Their Plausibility Four conditions must hold, and they are unevenly likely to. **The ground truth must be valid**: we inherit the NewsGuard scale together with its editorial judgements, and nothing internal to our data can test this. **The unit must be the publisher**, never the article: article-level accuracy is 0.60 and supports no claim about an individual piece. **Enough text must be available**. Small and recent outlets lack a rating most often, and they are the least likely to qualify. **Use must remain advisory**, ordering a review queue rather than rating anything autonomously. Our evidence covers English-language political news alone.

Negative Externalities and Responsible Deployment The principal risk is measured rather than hypothesised. Misclassification rates vary systematically with

political leaning (Section 5.4): they reach 59–72% for far-left and far-right outlets, against 30–34% for slightly right and non-labelled ones. The models also overestimate reliability at either extreme and underestimate it for slightly left-leaning outlets. Such error is systematic and reproducible, which distinguishes it in kind from that of an individual rater. Two further risks follow from presentation and use. A single level out of five reads as a measurement, although the l_4/l_5 boundary is close to unresolvable (Section 5.3). An inexpensive score is also a target, and surface features are easy to optimise. We therefore regard the following as necessary rather than advisable: publishing per-leaning error rates; abstaining when few articles are available; reporting an interval rather than a single level; and confining the output to ordering human review. The bias audit of Section 5.4 is constitutive of the contribution, not supplementary to it.

7 Discussion and Conclusion

In this study, we investigated whether the reliability of news publishers could be inferred from their articles, framing the task as a five-level ordinal classification problem aligned with NewsGuard ratings and evaluating four transformer-based models under a strict publisher-disjoint protocol.

Our findings demonstrate that language models can capture meaningful and generalizable signals of publisher reliability. Interestingly, the substantially larger Mistral 24B does not outperform smaller architectures, suggesting that performance is limited by dataset characteristics rather than model size. Despite operating with a 1,024-token context window, four times larger than the limit used for the other models, Mistral shows no measurable gain, which may indicate that the cues discriminating reliable publishers are concentrated in the beginning of articles and further supports the hypothesis that increasing model size or input length alone cannot overcome limitations imposed by weak supervision. A key finding is that publisher reliability should not be modelled as a property of isolated articles: it emerges as a collective signal that becomes more stable when evidence is aggregated across a source, which supports a shift from article-centric formulations toward source-level inference.

Future research should extend this framework to multilingual settings, explore confidence-aware aggregation strategies, and develop training objectives tailored to ordinal classification. The value of such systems, however, lies less in accuracy than in reach. What matters is whether they bring scarce editorial scrutiny to the outlets that currently go unrated, and whether they do so without penalising one part of the spectrum more than another.

Disclosure of Interests. The authors have no competing interests to declare.

References

1. Adji, S., de Lhoneux, M., Spinde, T.: HT-MAGPIE: a hierarchical transformer for article reliability classification. Datenbank-Spektrum **26**(1), 15–24 (2026). https:

//doi.org/10.1007/S13222-026-00526-Y, https://doi.org/10.1007/s13222-026-00526-y
2. Anonymous: Supplementary material for 'source-level reliability estimation under publisher-disjoint evaluation' (2026), https://anonymous.4open.science/r/supplementary-material-DD54/README.md, anonymised supplementary material accompanying this submission
3. Azizov, D., Mujahid, Z.M., AlQuabeh, H., Nakov, P., Liang, S.: SAFARI: Cross-lingual bias and factuality detection in news media and news articles. In: Findings of the Association for Computational Linguistics: EMNLP 2024. pp. 12217–12231. Association for Computational Linguistics (Nov 2024). https://doi.org/10.18653/v1/2024.findings-emnlp.712
4. Bazmi, P., Asadpour, M., Shakery, A.: Multi-view co-attention network for fake news detection by modeling topic-specific user and news source credibility. Information Processing and Management **60**(1) (2023)
5. Bohacek, M., Bravansky, M., Trhlík, F., Moravec, V.: Czech-ing the news: Article trustworthiness dataset for Czech. In: Barnes, J., De Clercq, O., Klinger, R. (eds.) Proceedings of the 13th Workshop on Computational Approaches to Subjectivity, Sentiment, & Social Media Analysis. pp. 96–109. Association for Computational Linguistics, Toronto, Canada (Jul 2023). https://doi.org/10.18653/v1/2023.wassa-1.10, https://aclanthology.org/2023.wassa-1.10/
6. Burdisso, S., Sanchez-cortes, D., Villatoro-tello, E., Motlicek, P.: Reliability estimation of news media sources: Birds of a feather flock together. In: Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers). pp. 6900–6918. Association for Computational Linguistics (2024). https://doi.org/10.18653/v1/2024.naacl-long.383
7. Devlin, J., Chang, M.W., Lee, K., Toutanova, K.: Bert: Pre-training of deep bidirectional transformers for language understanding. In: Proceedings of the 2019 conference of the North American chapter of the association for computational linguistics: human language technologies, volume 1 (long and short papers). pp. 4171–4186 (2019)
8. Fogg, B., Tseng, H.: The elements of computer credibility. In: Proceedings of the SIGCHI Conference on Human Factors in Computing Systems. p. 8087. CHI '99, Association for Computing Machinery (1999). https://doi.org/10.1145/302979.303001
9. Kim, A., Moravec, P., Dennis, A.: Combating fake news on social media with source ratings: The effects of user and expert reputation ratings. Journal of Management Information Systems **36**(3), 931–968 (2019). https://doi.org/10.1080/07421222.2019.1628921
10. Leite, J., Razuvayevskaya, O., Bontcheva, K.e.a.: Weakly supervised veracity classification with llm-predicted credibility signals. EPJ Data Sci. **14**(15) (2025)
11. Lin, H., Lasser, J., Lewandowsky, S., Cole, R., Gully, A., Rand, D.G., Pennycook, G.: High level of correspondence across different news domain quality rating sets. PNAS nexus **2**(9), pgad286 (2023)
12. Liu, Y., Ott, M., Goyal, N., Du, J., Joshi, M., Chen, D., Levy, O., Lewis, M., Zettlemoyer, L., Stoyanov, V.: Roberta: A robustly optimized bert pretraining approach. In: International Conference on Learning Representations (2020), https://openreview.net/forum?id=SyxS0T4tvS
13. Longo, A., Zappatore, M., Colella, D.D., Zampino, F., Calò, A.: xfair: A multi-layer approach to data fairness assessment and data fairification. ACM J. Data

Inf. Qual. **17**(4), 25:1–25:29 (2025). https://doi.org/10.1145/3769113, https://doi.org/10.1145/3769113

14. Mauricio, P.G., Quesada-Lopez, C., Martinez, A., Jenkins, M.: Automatic classification of web news: A systematic mapping study. In: Intelligent Systems and Applications (IntelliSys 2020). pp. 558–574. Springer (2020). https://doi.org/10.1007/978-3-030-52249-0_44
15. Metzger, M., AJ, F.: Digital Media, Youth, and Credibility. MIT Press (2007 (1st edition))
16. Mistral AI: Mistral small 3. Model card: `mistralai/Mistral-Small-24B-Base-2501` (2025), https://mistral.ai/news/mistral-small-3/, 24B-parameter dense model
17. Monahov, A.: Improved accuracy metrics for classification with imbalanced data and where distance from the truth matters, with the wconf r package. Available at SSRN 4802336 (2024)
18. Nakov, P., An, J., Kwak, H., Manzoor, M.A., Mujahid, Z.M., Sencar, H.T.: A survey on predicting the factuality and the bias of news media. In: Findings of the Association for Computational Linguistics: ACL 2024. p. 1594715962 (2024). https://doi.org/10.18653/v1/2024.findings-acl.944
19. Pratelli, M., Petrocchi, M.: A structured analysis of journalistic evaluations for news source reliability. In: Workshop Proceedings of the 16th International AAAI Conference on Web and Social Media (2022), https://doi.org/10.36190/2022.51
20. Pratelli, M., Saracco, F., Petrocchi, M.: TROPIC - trustworthiness rating of online publishers through online interactions calculation. In: Hauff, C., Macdonald, C., Jannach, D., Kazai, G., Nardini, F.M., Pinelli, F., Silvestri, F., Tonellotto, N. (eds.) Advances in Information Retrieval - 47th European Conference on Information Retrieval, ECIR 2025, Lucca, Italy, April 6-10, 2025, Proceedings, Part IV. Lecture Notes in Computer Science, vol. 15575, pp. 407–412. Springer (2025). https://doi.org/10.1007/978-3-031-88717-8_30, https://doi.org/10.1007/978-3-031-88717-8_30
21. Przybyla, P., Soto, A.: When classification accuracy is not enough: Explaining news credibility assessment. Information Processing and Management **58**(5), 102653 (2021). https://doi.org/https://doi.org/10.1016/j.ipm.2021.102653, https://www.sciencedirect.com/science/article/pii/S0306457321001412
22. Rogers, A., Kovaleva, O., Rumshisky, A.: A primer in bertology: What we know about how bert works. Transactions of the association for computational linguistics **8**, 842–866 (2021)
23. Spinde, T., Richter, E., Wessel, M., Kulshrestha, J., Donnay, K.: What do twitter comments tell about news article bias? assessing the impact of news article bias on its perception on twitter. Online Soc. Networks Media **37-38**, 100264 (2023). https://doi.org/10.1016/J.OSNEM.2023.100264, https://doi.org/10.1016/j.osnem.2023.100264
24. Srba, I., et al.: A survey on automatic credibility assessment using textual credibility signals in the era of large language models. ACM Trans. Intell. Syst. Technol. (Sep 2025). https://doi.org/10.1145/3770077, just Accepted
25. Team, L.: The llama 3 herd of models. CoRR **abs/2407.21783** (2024). https://doi.org/10.48550/ARXIV.2407.21783, https://doi.org/10.48550/arXiv.2407.21783
26. Verhoeven, I., Mishra, P., Shutova, E.: Yesterdays news: Benchmarking multi-dimensional out-of-distribution generalization of misinformation detection models. Computational Linguistics pp. 1–58 (2025)
27. Wardle, C., Derakhshan, H.: Information disorder: Toward an interdisciplinary framework for research and policy making. Tech. rep., Council of Europe (2017)

28. Williams, E.M., Carragher, P., Carley, K.M.: Bridging social media and search engines: Dredge words and the detection of unreliable domains. In: An, J., Lin, Y., Mejova, Y., Mustafaraj, E., Kulshrestha, J., Weber, I. (eds.) Proceedings of the Nineteenth International AAAI Conference on Web and Social Media, June 23-26, 2025, Copenhagen, Denmark. pp. 2030–2043. AAAI Press (2025). https://doi.org/10.1609/ICWSM.V19I1.35917, https://doi.org/10.1609/icwsm.v19i1.35917
29. Yang, K.C., Menczer, F.: Accuracy and political bias of news source credibility ratings by large language models. In: Proceedings of the 17th ACM Web Science Conference 2025. p. 127137. Websci '25, Association for Computing Machinery, New York, NY, USA (2025). https://doi.org/10.1145/3717867.3717903, https://doi.org/10.1145/3717867.3717903